

Segmented Changelog

Jun Wu

Facebook

April 15, 2019

Goals

Improve $O(\text{changelog}^\dagger)$ DAG operations.

- ▶ $O(1)$ -ish clone with edenfs
- ▶ $O(\log \text{changelog})$ -ish ancestor related DAG calculations

Improve $O(\text{changelog})$ set operations.

- ▶ $a + b$
- ▶ $a - b$
- ▶ $a \& b$

[†]Number of commits in the repo

Non-goals

Filter operations. Expect they run on a small subset.

- ▶ `author(alice)`
- ▶ `date(2017)`
- ▶ `branchpoint()`
- ▶ `merge()`

File DAG operations. Expect server-side support (ex. fastlog).


- ▶ `follow(path, commit)`

Commit Identities



[grapes, raspberry, watermelon, orange, kiwi]

$O(N)$ space




1:5

$O(1)$ space

Segments

Segments and parents

A Segment contains a list of numbers.
Numbers are sorted topologically in DAG.



Parents of a segment: `parents(segment)` - segment. Stored with the segment.

Flat Segments

Constraints:

- ▶ $x:y$ (using revision numbers)
- ▶ $\text{heads}(x:y)$ is y
- ▶ $\text{roots}(x:y)$ is x
- ▶ $((x:y) - x) \ \& \ \text{merge}()$ is empty




Properties:

- ▶ Reduce DAG complexity to $O(\text{merges})$
- ▶ Parents of x are segment parents.
- ▶ Parent information is loseless.

Flat Segments


Example



High-Level Segments

Constraints:

- ▶ $x:y$ (using revision numbers)
- ▶ $\text{heads}(x:y)$ is y



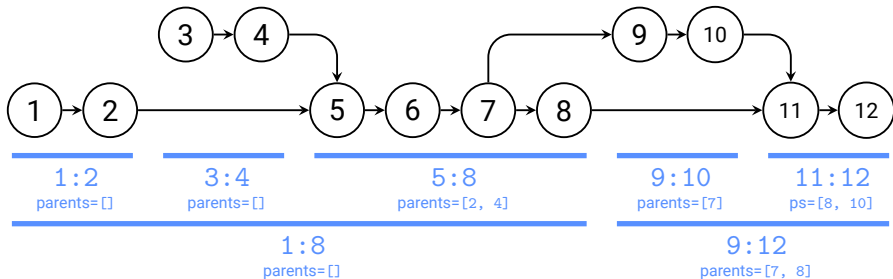
Properties:

- ▶ Compress commits across merges
- ▶ Parent information is lossy.[†]

[†]Cannot get parents of arbitrary commit by high-level segments only.

High-Level Segments

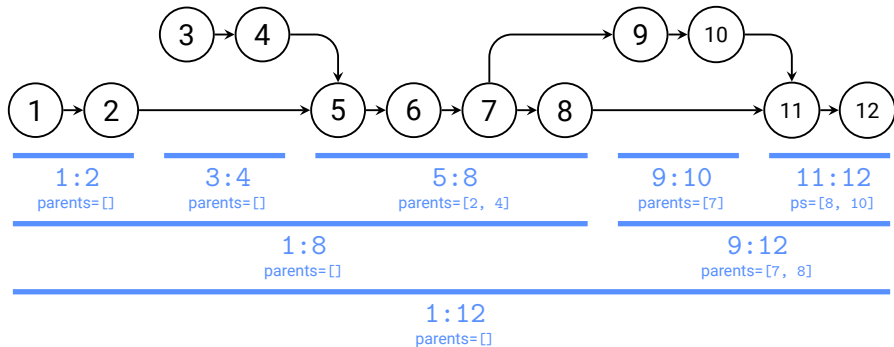
Example



Segment Size = 3. A high-level segment contains 3 lower-level segments at most.

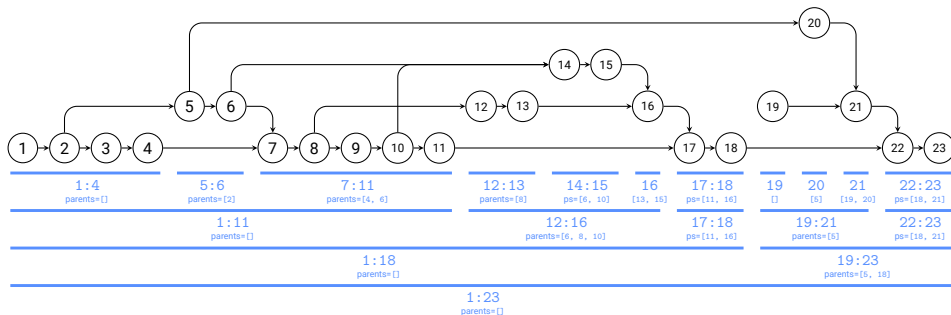
High-Level Segments

Example



High-Level Segments

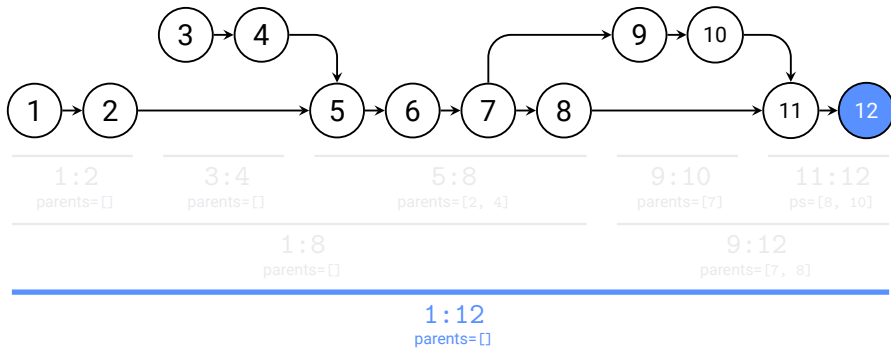
Example



High-Level Segments

Example: Ancestors Selection


Selecting ::12



High-Level Segments

Example: Ancestors Selection

Selecting $::11$

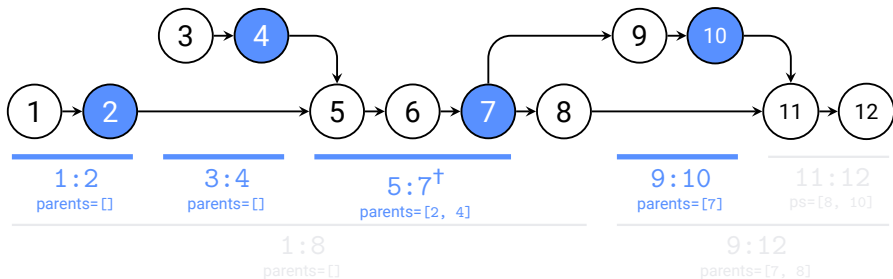


[†]11:11 is generated from 11:12.

High-Level Segments

Example: Ancestors Selection

Selecting $::10$




[†]5:7 is generated from 5:8.

High-Level Segments

Example: Ancestors Selection

Selecting ::8



High-Level Segments

Example: Common Ancestor

```
ancestor(10, 8)
= max(::10† & ::8)
= max((1:7+9:10) & 1:8)
= max(1:7)
= 7
```





[†] ::10 can be lazy[‡] to avoid 1:2 and 3:4 lookups.

[‡] Laziness has a cost. Non-lazy version has $O(\log \text{ changelog})$ -ish overhead.

Assigning Numbers


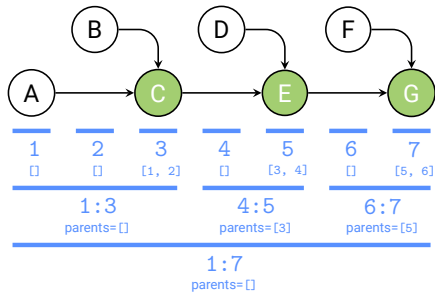
Flat Segments



Use Depth-First Search, not Breadth-First Search, from a merge.

Assigning Numbers

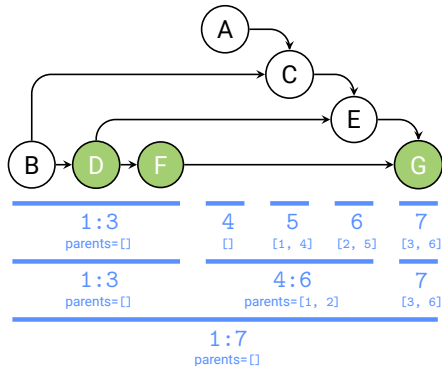
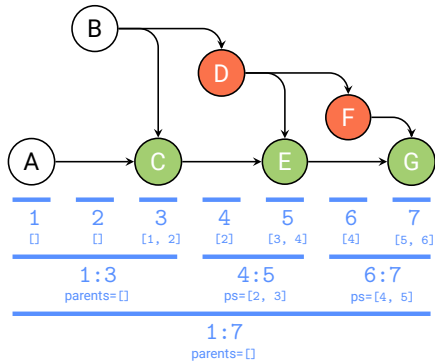
Merges



Use Depth-First Search.

Assigning Numbers

Merges



Pick parent branch with less merges first.

Assigning Numbers

Algorithm

To assign a number for x , check its parents.

- ▶ If all parents have numbers, assign the next available number to x .
- ▶ Otherwise, pick the parent branch with less merges. Assign it recursively.

Real-world Repos

Repo	Commits	Flat	Level 2 [†]	Level 3	Level 4
fbsource	millions	19961	1325	87	6
mozilla	469k	34179	2390	149	8
cpython	98k	21744	1367	81	4
pypy	74k	9008	609	39	2
git	55k	23884	1534	90	5
hg	42k	4654	297	17	0

[†]Segment Size = 16. Last segment per level is removed.

New Structures

- ▶ Number - Commit Hash Mapping
 - ▶ Large
 - ▶ Stored Server-side
- ▶ Commit Hash - Commit Data (user, date, message) Mapping
 - ▶ Larger
 - ▶ Stored Server-side
- ▶ Segments
 - ▶ Tiny (<1MB)
 - ▶ Calculated Server-side
 - ▶ Stored client-side